
PRODUCT EXPERIENCE AUDIT

MOSAIX

Where workforces are designed for the future.

A read of the MOSAIX skills-intelligence platform, from five different angles, grounded in the live demo, and written to find the redesign that turns capability into trust.

PREPARED BY	SUBJECT	DATE	STATUS
IDYeah Studio	MOSAIX · Skills Intelligence	July 2026	Confidential

CONTENTS

What's inside

01	Executive summary	<i>The short version</i>
02	How I read it	<i>Five angles, one evidence base</i>
03	Assumptions	<i>What I looked at, and what I didn't</i>
04	What's genuinely working	<i>The core worth protecting</i>
05	The central finding	<i>Coherence, not capability</i>
06	The five angles	<i>Narrative · Experience · Craft · Investor · Buyer</i>
07	One more read	<i>The pitch deck</i>
08	Findings at a glance	<i>Severity matrix</i>
09	The redesign roadmap	<i>Sequenced by return</i>
10	The opportunity	<i>Why this is the moment</i>
11	Evidence & sources	<i>What this is grounded in</i>

The short version

MOSAIX has already built the hard part. The redesign’s job is to make the product sound like what it already is.

CAPABILITY	COHERENCE	THE VERDICT
Strong	At risk	Redesign, not rebuild

Here is the short version. MOSAIX has already built the hard part. The succession, analytics, talent-mapping and org-design screens are the real thing. A talent leader could take them into a board review and hold their own. That is rare, and it is what the rest of this note builds on.

The gap is not capability. It is **coherence**. A strong core is wrapped in a shell that reads as unfinished, and in a few places it argues with itself. Employee and position counts do not match from one screen to the next. A meme-and-badge layer sits one click from enterprise analytics. “Apply Coupon Code” shows up on a product sold on ISO 27001 and SOC 2 trust. Four modules say “Coming Soon” right there in the main navigation. And a single 25-item sidebar is shown, in full, to everyone.

This matters more for MOSAIX than for most products. Eightfold owns the data. Gloat owns mobility. Workday owns the install base. MOSAIX has neither a data moat nor a distribution moat. What it can own is the experience: broad, fast, and trustworthy. So here **the experience is the moat**. Every mismatch and dead end chips at the one thing the company can actually defend, and the global launch will amplify it, not hide it.

I looked at this from five angles and they all landed in the same place, which is the good news. This is a **coherence redesign, not a rebuild**. The capability is already there. The job is to close the gap between the two. That is the exact thing this category rewards.

Five angles, one evidence base

I do not begin in a design tool. I begin by understanding the brand, the market, and the person the product is really for, and then I design.

I went through the live MOSAIX demo screen by screen and read every claim against the deck and against how this category is won in 2026. To stop any one instinct from running the show, I read the product from five different angles, so a finding only stays in if it holds up as narrative, as experience, as craft, as investment risk, and from the buyer’s own chair. Every finding points at a real screen. None of it is guesswork.

THE FIVE ANGLES

01	Positioning & Narrative	<i>Does the experience prove the story the deck tells?</i>
02	Experience & Architecture	<i>The path from first click to first value.</i>
03	Visual System & Craft	<i>One system, one voice, enterprise craft.</i>
04	Investment & Growth	<i>Does the product underwrite the growth story?</i>
05	The Buyer’s Chair	<i>Seen from the seat that has to sign off.</i>

Grounded in the MOSAIX US deck, 24 screenshots of the live demo (Org Admin session, sample data only), and 2026 talent-intelligence and enterprise-HR buyer research.

What I looked at, and what I didn't

Before the findings, a few ground rules, so you can weigh them fairly.

I looked at the **demo instance**. There is no live customer data in it. The people, the counts and the assessment cycles are all sample data, seeded for a walkthrough. So when I point at numbers that do not reconcile, I am not claiming the production database is broken. I am saying the demo a prospect gets shown does not hold together, and the demo is where the first impression is made.

Some of the product is not built yet, and that is fine. **Four modules openly say “Coming Soon.”** I have treated those as roadmap, not as failures. Where a screen is a placeholder for something planned, I have not marked it as a defect. What I have flagged is the choice to leave those placeholders sitting in the main navigation while a buyer is clicking around.

I reviewed what I could reach as an Org Admin in a single sitting, which came to **twenty-four screens**. I did not have live integrations, real assessment data, or a customer's actual org loaded. If any of that changes the picture, tell me and I will adjust.

So I have kept two things separate: “this is unfinished or seeded,” and “this choice would cost you even after it is finished.” Almost everything I recommend is about the second one.

The core worth protecting

Before I pick at anything, let me be precise about what MOSAIX gets right. The redesign protects these. It does not touch them.

- **Succession Planning is best in class.** One clear decision surface. The framing stats up top, then readiness columns (Ready Now, 12, 24, 24+ months) that match how a talent leader already thinks. The Org and 9-Box toggle and “Modify Successors” let you act in the same place you diagnose.
- **Succession Analytics is genuinely decision-grade.** Bench depth, performance and potential misalignment, and the one most tools skip: top performers and high-potentials who are being missed. Those are the questions a board actually asks.
- **Talent Mapper is a tool people would really use.** Honest, dense search across department, skill, performance, location and tenure, over employee cards with skill chips and a clear next step.
- **Org Design Engine shows nerve.** A “describe your org” prompt paired with a plain “AI-generated, indicative only” note. That is a trust instinct most AI features skip.
- **The visual base is strong.** Clean components, one consistent card pattern, warm pastel icon tiles, and real discipline on the dense screens.
- **The onboarding instinct is right.** Dropping every new account into a populated sample org is the correct call. It just needs to reach every screen.

PROTECT & AMPLIFY

Decision-grade core

- Succession Planning & Analytics
- Talent Mapper search + skill chips
- Org Design Engine (AI + honest note)
- Clean, consistent visual base
- Populated sample-org onboarding instinct

FIX TO UNLOCK IT

Coherence, not capability

- Cross-screen data that doesn't reconcile
- Two brand voices in one skin
- SMB chrome on an enterprise product
- 25-item nav, no role-tailoring
- “Coming Soon” dead-ends in the main nav

Coherence, not capability

Across all five angles, narrative, experience, craft, investment, and the buyer's own chair, I kept arriving at the same sentence:

MOSAIX's problem is not what it can do. It is that the experience does not yet agree with itself.

Three fault lines produce almost every finding in this note.

- 1 Trust cracks in the data.** The welcome flow seeds 100 sample employees. Talent Mapper shows 788. Succession Planning says 5 positions and 3 critical roles. Succession Analytics says 150 and 24. This is a demo, so these are seed numbers, not a broken database. But the buyer does not know that. For a product that sells "real-time skills intelligence," numbers that do not line up from one screen to the next quietly tell them the intelligence cannot be trusted. Fix the demo, because the demo is what they see.
- 2 Two products wear one skin.** Enterprise analytics sits a click away from Meme of the Day, Riddle of the Day, and a wall of locked badges ("Selfie Starter," "Wordsmith Wizard"). The voice splits in half, and the enterprise buyer resolves the split against you.
- 3 The shell reads early, or SMB.** "Apply Coupon Code" and a trial badge on an ISO and SOC 2 product. Four "Coming Soon" dead ends in the main nav. A 25-item sidebar with no role-tailoring. Naming that drifts. The same primary button repeated on one screen.

None of this touches the capable core. That is the whole opportunity.

Five ways of seeing one product

I read each angle on its own terms, and asked each to lead with what works before going deep on what holds the product back. The emphases differ. The conclusion does not.

Each angle opens with where it stands and the read in one line, then the evidence underneath.

01 Positioning & Narrative

WHERE IT STANDS Strong story, undercut by the shell

IN ONE LINE The pitch is proven on three screens, then argued with on the rest. Close that gap and the story sells itself.

What's working

The story-to-product translation is strongest exactly where it needs to be. Succession Planning and Analytics are credible and board-ready: readiness columns, coverage, a 9-Box toggle, department bars, export. This is the hardest thing in the category to build, and it is the best thing in the product. It proves the “real-time skills intelligence” line instead of just claiming it. Talent Mapper backs it up with real search depth (“Showing 20 of 788”). Org Design Engine is the differentiated bet, and its honest “indicative only” note shows a team that respects a skeptical buyer. Against Eightfold’s data and Workday’s incumbency, that trio (breadth, craft and decision-readiness) is the only wedge MOSAIX can actually win on.

Where the experience argues with the story

The problem is not capability. It is **coherence**, and here coherence is the differentiator, not a coat of paint.

Two products, one skin. Enterprise analytics sits one click from Meme of the Day (yes, the “Woman Yelling at Cat” meme), a Riddle of the Day, and a wall of locked badges. A buyer mid-evaluation cannot square that with an ISO 27001 and SOC 2 pitch, so they square it against you.

Trust signals that fight themselves. “Apply Coupon Code” and a TRIAL badge sit in the top bar of a product selling enterprise certification. Coupons say SMB self-serve. That undercuts the procurement-led motion the succession screens imply.

The launchpad undersells the product. Home greets you with 2 tiles against a 25-item nav, then hands you four “Coming Soon” dead ends and a run of empty screens. So the opening shows the least of what MOSAIX does.

Numbers that do not line up. 100 sample employees, then 788, then 5 positions, then 150. In a demo built to sell “real-time intelligence,” that is the most corrosive signal in here.

Naming that drifts. Mentor Marketplace becomes “Mentorverse.” My Career Planner becomes My Career Arc. A “9 Blocker” banner that should say 9-Box. Small on their own. Together they whisper “unfinished.”

The moves I would make first

- 1 **Pick one voice, and make it the enterprise one.** Put the fun stuff behind a toggle HR controls. Lead every path with succession, analytics and org design.
- 2 **Rebuild the first five minutes as a proof tour** into the three strong screens, with the sample data reconciled.

- 3 **Strip the SMB tells and the dead ends.** Lose the coupon chrome, quiet the trial badge, move “Coming Soon” into a real roadmap view.
- 4 **Make the wedge obvious and manager-first,** and settle on one name per thing so the navigation itself reads like a brand.

O2 Experience & Architecture

WHERE IT STANDS The hard part is built, then buried

IN ONE LINE The capability is real. Navigation, first-run state and cross-screen trust are where it leaks. The fix is coherence, not reconstruction.

MOSAIX does not lack capability. It just surfaces the capability it has unevenly. The strongest screens are among the hardest things to design in this category, and they are designed well. So this is about coherence, not reconstruction.

What's working

Four screens are genuinely review-ready and should anchor the redesign. **Succession Planning** is best in class: the stats frame the situation, the readiness columns match how a reviewer thinks, and the Org and 9-Box toggle lets you act where you diagnose. **Succession Analytics** carries that into aggregate insight. **Talent Mapper** is an honest, usable search. **Org Design Engine** shows nerve with its “indicative only” note. The pastel icons, the single card pattern, and the populated-sample-org instinct are all right.

Findings

One note before I start: this is a demo with seeded data, so some empty screens and mismatched counts are almost certainly sample artifacts. I kept them in because this is the exact build a prospect gets shown.

F1. One undifferentiated 25-item sidebar. Critical. A single Org Admin sees all ~25 destinations at once. The bands are labelled but not gated by role. That breaks role-based surfacing and progressive disclosure. **Direction:** role-resolved navigation.

F2. Onboarding asks for precision before it earns trust. High. A required domain field, heavy free text at step one, and overlapping region chips (Americas already contains North America and Latin America; EMEA already contains Europe, the Middle East and Africa). **Direction:** a clean, non-overlapping region list, an optional domain, and defer the strategy questions.

F3. The launchpad undersells the product. High. 2 tiles against a 25-item nav. **Direction:** a role-aware launchpad that surfaces the best next actions.

F4. Empty screens even with a sample org loaded. High. Career Arc, Career Recommendations, Project Marketplace, My Team and Talent Metrics all render empty. **Direction:** seed every first-run screen, and treat empty states as guidance.

F5. Over-built controls on empty screens. Medium. The empty IDP ships five view toggles. **Direction:** hide controls until there is content.

F6. The same primary button, repeated. Medium. “Create New Session” three times on one screen. “Generate” twice. **Direction:** one primary per screen.

F7. Counts that do not reconcile across screens. Critical. 100 vs 788 employees. 5 vs 150 positions. 3 vs 24 critical roles. In a product that sells intelligence, this reads as a trust failure, seed data or not. **Direction:** one source of truth per number, and make cross-screen agreement a release gate on the demo.

F8. Naming drift. Medium. Mentorverse vs Mentor Marketplace. “9 Blocker” vs 9-Box. **Direction:** one agreed name per concept.

F9. Managers have to work too hard. High. The manager view sits inside admin-scale navigation with no fast path to the two things a manager does: review a report and log feedback. **Direction:** a dedicated, low-chrome manager mode.

What I would fix first

- 1 Reconcile the numbers (F7). Cheapest fix, biggest trust return.
- 2 Role-resolved nav (F1) and a role-aware launchpad (F3).
- 3 A frictionless manager mode (F9).
- 4 First-run integrity: seed coverage and guided empty states (F4, F5).
- 5 Onboarding taxonomy (F2) and a consistency pass (F6, F8).

03 Visual System & Craft

WHERE IT STANDS A capable base speaking in two voices

IN ONE LINE A modern base carrying two brand voices and a few ad-hoc primitives. One token layer and one voice away from enterprise-grade.

What's working

The visual base is genuinely good. A clean, current component language, one card pattern that repeats predictably, and a warm pastel icon-tile system that makes dense action grids easy to scan. The discipline on Succession Planning and Analytics is real: heavy density, no visual panic. That is a system you can push, not one you throw away.

Findings

F1. System and tokens. Three different treatments carry “this is the action”: gradient headers and buttons, a flat purple “Add,” and blue text links. Three weights for one idea means there is no single definition of primary, secondary and link. That is a token gap, not a taste gap.

F2. Type and hierarchy on dense screens. Hierarchy leans on the card edges instead of a real heading, label and value scale, so the eye does not land on the numbers that matter first. These screens need a proper type ramp with tabular figures.

F3. The tonal split. The product speaks two visual languages: enterprise analytics, and consumer gamification (the meme, the badges called “Selfie Starter” and “Wordsmith Wizard”). One product should not talk in two voices.

F4. Craft slips. Broken, grayed article thumbnails on the first screen. The “Coming Soon” rocket reused across four modules. The little info icon present on some cards, missing on others. Small things that add up to “unfinished.”

F5. Enterprise fit. The visual language signals SMB self-serve where it should signal enterprise. “Apply Coupon Code” and a “TRIAL VERSION” badge are shop chrome on an enterprise surface.

F6. Accessibility. Light text on the gradient headers, badge state shown by fade alone, and readiness carried only by color. Each needs a second, non-color signal.

F7. An icon for everything. Nearly every nav item, module and feature card carries its own bespoke icon, so the product ships 25-plus little symbols a user is asked to learn, and none are distinctive enough to remember. Past a point, icons stop aiding recognition and just add noise. They should earn their place, not decorate every row.

What I would do

- 1 Put in a real token layer that collapses the gradient, flat-purple and blue-link into one governed set (F1).

- 2 Settle on one voice, and give the fun layer its own opt-in home (F3).
- 3 Ship a type ramp for density, with tabular figures on the analytics screens (F2).
- 4 Fix the craft floor: one shared empty and coming-soon component, repaired thumbnails, a consistent info affordance (F4).
- 5 Swap the shop chrome for enterprise trust cues, contrast-test the gradients, and add a second signal to anything shown by color (F5, F6).
- 6 Tame the icon set: a small, consistent library with clear metaphors, used for wayfinding and status, not sprinkled on every card (F7).

04 Investment & Growth

WHERE IT STANDS Fundable core, discounting its own story

IN ONE LINE A real, category-serious product whose experience is discounting the growth story. Here the experience is the only moat.

On paper this looks like a company that has found early traction: 267% growth in 2025, six paying logos in the 1k–12k employee band, 25k-plus headcount under management, ISO 27001 and SOC 2 done, a global launch set for Q2 2026. Log in, and the product does not embarrass the deck. The question is narrower: does the live experience back the story, or quietly discount it?

Strengths. The hardest things are the best built. Succession Planning and Analytics are the screen a talent leader actually runs a review from. Talent Mapper is a real, usable search. Org Design Engine pairs a generative prompt with an honest disclaimer, which reads as maturity, not AI theater. Breadth is the bet: a full Define, Assess, Build, Plan, Grow lifecycle, standalone, is a coherent wedge against Workday's suite and Gloat's narrower focus. The team can build, which takes real execution risk off the table.

Flags a diligence process would catch. The big one is the counts that do not reconcile across screens. I know it is a demo with seed data. But it is the demo a prospect gets shown, and for a product selling "real-time skills intelligence," numbers that do not tie out are exactly the thing a sharp buyer catches live. Second, four "Coming Soon" modules sit in the main nav, advertising what is not built. Third, one role sees all 25 destinations, and adoption, not capability, is the game. Fourth, two products in one skin, plus coupon and trial chrome that fights the ISO and SOC 2 posture.

The compounding fact: no data moat, no distribution moat. Eightfold has the data, Gloat the use case, Workday the install base. What is left to own is the experience. That is the only moat here, and every mismatch chips at it.

Why design is the highest-return lever right now. In 2026, AI alone does not sell. Buyers weigh the everyday UX and whether managers actually adopt it, and they want proof the AI changed a real decision. MOSAIX's capability clears that bar. Its coherence does not yet. A focused redesign moves the three numbers that matter: win rate (reconciled data removes the demo moment that loses deals), time to value (role-tailored entry), and retention (flows people actually adopt). It is cheaper than more engineering, and it is the thing to fix before a global launch makes the first impression permanent.

05 The Buyer's Chair

WHERE IT STANDS An impressive engine, not yet shortlist-ready

IN ONE LINE You've built the succession engine I'd take to my board. Make every number agree and lose the coupon-code costume, and you're on my shortlist.

This last angle is not mine. I wrote it from the chair of the person who actually signs the cheque: a senior HR leader at a large company that already runs Workday. Here is how the demo lands for them.

Fifteen minutes. Same as every "AI for HR" pitch that hits my desk. I run People Operations for about 12,000 people, Workday is my system of record, and my name is on every succession slate the board sees. I am not grading ambition. I am grading whether I would stake my credibility on a number this puts in front of my CEO.

Where you earn my confidence

The succession suite is the real thing, and it is not close. **Succession Planning** is the screen I would take straight into a talent review: the readiness bands, the coverage, the critical-role and no-successor counts, the 9-Box toggle, "Modify Successors" right where I would reach for it. **Succession Analytics** goes further than most tools dare: bench depth, misalignment, and the top performers and high-potentials I am missing, which is exactly what my board asks and my current stack answers slowly. **Talent Mapper** is something a manager would actually use to build a team. Whoever built these gets the job. I will also give you credit for shipping an AI estimate with an honest "indicative only" note.

Where you lose me

Then I look up and the enterprise credibility drains out of the room. "**Apply Coupon Code**" sits in my top nav, next to a "**TRIAL VERSION**" badge. I am weighing a six-figure platform and you are showing me the checkout furniture of a \$12-a-month app. When the coupon and the compliance badge disagree, my board believes the coupon.

Then the feed: a meme, a riddle, badges called "Selfie Starter" and "Wordsmith Wizard," with **broken images on the very first screen**. It reads as condescending to my leaders and unserious to my board. And because I am an Org Admin you show me *everything at once*, all 25 destinations, four "Coming Soon" dead ends in the main nav. You never asked who I am, so you never tailored what I see.

Here is the one that ends it. Your welcome screen seeds **100 employees**. Talent Mapper says **788**. Succession Planning says **5 positions and 3 critical**. Succession Analytics says **150 and 24**. Four screens, four different headcounts, in a product that sells "real-time intelligence." My whole job is *defensible* people-data. If your own demo cannot keep its headcount straight, I cannot trust a single number you would feed my succession review. I know it is only sample data, and that is exactly why it stings: it was avoidable.

What would move me to a shortlist

- 1 **Reconcile the data. Once, everywhere.** Until the numbers agree, nothing else matters.
- 2 **Give me a role-aware view.** A CHRO cockpit first. The gamification opt-in and segmented. Get “Coming Soon” out of my main nav.
- 3 **Dress it like enterprise software.** Retire the coupon and trial chrome from the leader view. Put ISO, SOC 2 and the Workday integration where I will see them in the first five minutes.
- 4 **Prove the AI changed a real decision.** Show me one live case where a recommendation moved a manager to name a successor.

The pitch deck

The five angles read the product. This one reads the deck, because the story you tell about MOSAIX is its own product, and it is the other thing I would want to redesign.

WHERE IT STANDS

Strong facts, argued before they are earned

IN ONE LINE

The proof points are genuinely good. The trouble is order and belief: the deck opens on credentials and features instead of a story, and asks the room to accept big numbers it has not yet shown.

What's working

- **The problem is well armed.** The WEF, McKinsey and Gartner stats (44% of core skills change by 2027, 63% of employers name skill gaps as the top barrier, 50% of leaders say they do not use the skills they already have) give the deck urgency and outside credibility.
- **The cost of inaction lands.** Naming missed revenue, rising cost, slower execution, attrition and execution risk turns a vague worry into a bill.
- **The five-step spine is memorable.** Define, Assess, Build, Plan, Grow gives the platform a shape a buyer can hold in their head.
- **The team is a real asset.** A century of combined HR leadership, an ex-Amazon product leader and a genuine CTO close a lot of doubt.
- **There are bold outcomes and real traction.** 40% lower hiring cost, roughly 30% more productivity, 3x faster transformation, six paying customers, 25k-plus headcount, a Q2 2026 launch.

Where the story loses power

- **It opens on the org chart, not the stakes.** Leading with the leadership team asks the room to care who you are before it cares about the problem. Open on the tension instead.
- **No named enemy, no wedge.** The deck never says why MOSAIX beats the suite (Workday) or the point tools (Gloat, Eightfold). Without that, the strong stats sell the category, not the company.
- **Big numbers, no proof.** 40% and 3x and 80% are exactly the claims a skeptical buyer discounts on sight, because nothing on the page stands behind them. Six customers are counted, none are shown.
- **Features interrupt the story.** The system-architecture slide is an operations detail dropped into the middle of a narrative. It answers a question no one has asked yet.
- **The close is flat.** The deck ends on "Thank you," and the traction slide reads as a row of unlabeled employee counts. No vision, no ask.

- **The human “why” never lands.** The mission gestures at people and their growth, but the body is stats, then features, then numbers. The feeling under the facts is missing.

Moves I would make

- 1 **Reorder to an arc:** hook, stakes, the shift (skills, not roles), why now and why you, how it works, one real customer before-and-after, then vision and a clear ask.
- 2 **Open on the shift, not the team.** Move the credentials to a later beat, where they close doubt instead of opening the show.
- 3 **Add one proof story.** Even anonymized, a single customer before-and-after does more than every percentage on the page.
- 4 **Name the wedge.** One slide on why MOSAIX wins where the suite is heavy and the point tools are narrow.
- 5 **Trade the architecture diagram** for a single “how it works” picture tied to an outcome.
- 6 **Close on a vision and an ask,** not a thank-you.

The severity matrix

Sixteen findings, consolidated across the five angles and ranked by the risk each poses to trust and adoption. The pattern worth noting: the Critical items are the cheapest to fix and the most damaging left alone.

2 CRITICAL

7 HIGH

7 MEDIUM

FINDING	SEEN FROM	SEVERITY
Cross-screen data does not reconcile (100 / 788 / 5 / 150 / 3 / 24)	EXPERIENCE · INVESTOR · BUYER	CRITICAL
Single 25-item sidebar; no role-based surfacing	EXPERIENCE · BUYER	CRITICAL
Launchpad (2 tiles) under-promises a 25-module product	EXPERIENCE · NARRATIVE	HIGH
Under-guided empty states despite a loaded sample org	EXPERIENCE	HIGH
Manager path buried; adoption task-cost too high	EXPERIENCE · INVESTOR	HIGH
Two brand voices: enterprise analytics vs consumer gamification	VISUAL · NARRATIVE · BUYER	HIGH
SMB self-serve chrome (coupon / trial) on an ISO-SOC2 product	NARRATIVE · VISUAL · BUYER	HIGH
Four “Coming Soon” modules live in primary navigation	NARRATIVE · INVESTOR · BUYER	HIGH
Onboarding: overlapping geo taxonomy + heavy free-text at step 1	EXPERIENCE	HIGH
Ad-hoc action primitives; no semantic token layer	VISUAL	MEDIUM
Duplicated primary CTAs (Org Design ×3; Career Recs ×2)	EXPERIENCE · VISUAL	MEDIUM
Naming drift (Mentorverse↔Mentor Marketplace; “9 Blocker”↔9-Box)	EXPERIENCE · NARRATIVE	MEDIUM
Over-engineered controls on empty screens (5 IDP toggles)	EXPERIENCE	MEDIUM

FINDING	SEEN FROM	SEVERITY
Craft defects (broken thumbnails; repeated rocket; (i) affordance)	VISUAL	MEDIUM
Icon overload: 25+ bespoke icons dilute recognition	VISUAL · EXPERIENCE	MEDIUM
Accessibility: gradient contrast + color-only status encoding	VISUAL	MEDIUM

Sequenced by return, not by size

This is a coherence redesign, not a rebuild, so the highest-value work is also among the fastest. I sequence trust first, because every later gain is undermined without it.

01 Trust foundation

DAYS-WEEKS

LOW EFFORT · HIGH IMPACT

Reconcile all cross-screen data to one source of truth; remove coupon and trial chrome from the enterprise view; pull “Coming Soon” modules out of primary nav; repair broken images; settle a governed lexicon.

PAYOFF Removes the demo-killing moments and restores baseline credibility.

02 Coherence redesign

WEEKS

MEDIUM EFFORT · HIGH IMPACT

Role-resolved navigation and a role-aware launchpad; a frictionless manager mode; one settled brand voice with the engagement layer given an opt-in home.

PAYOFF Lifts manager adoption and time to value, the category’s deciding criteria.

03 System & first-run

WEEKS

MEDIUM EFFORT · COMPOUNDING

A real design-token layer and a type ramp for dense screens; a single empty-state and seed-coverage system; onboarding taxonomy and sequencing; an accessibility pass.

PAYOFF Keeps the product coherent as it scales past the global launch.

04 Differentiation

ONGOING

ONGOING · STRATEGIC

Put the strongest screens front and center; show the proof that the AI changed a real decision.

PAYOFF Converts capability into the one thing 2026 buyers reward: trusted outcomes.

Why this is the moment

All five angles land in the same place. **MOSAIX has built the expensive, hard-to-copy part, and it is underselling it with cheap, fixable incoherence.** That is the best position a redesign can inherit. The question is not “can this product be good.” It is “will the market see how good it already is before the launch sets the first impression in stone.”

This is the work I do. I do not start in a design tool. I start by understanding the brand, the market, and the person the product is really for, and then I design. This note is that method applied before I redraw a single pixel: every finding tied to a real screen, weighed against how this category is won, and read from five different angles.

Here is what a focused engagement would deliver, in order.

- 1 **A trust foundation, fast.** Reconcile the demo data, retire the SMB chrome, get “Coming Soon” out of the main nav, and settle the naming. Cheapest work, biggest credibility return, done first.
- 2 **A coherence redesign.** Role-resolved navigation and a role-aware launchpad, a frictionless manager mode, and one settled voice that leads with the decision-grade core.
- 3 **A design system that holds.** A real token layer, a type ramp for the dense screens, one first-run system, and an accessibility pass, so the product stays coherent as it scales past launch.
- 4 **A differentiation story.** Make the strong screens the hero of the first five minutes, and surface the proof that the AI changed a real decision.

There are two things here worth redesigning, not one: the product, and the deck that sells it. The deck read a few pages back is a first pass at the second.

The invitation is simple. Let’s take a capable product, and the story you tell about it, and make both obviously trustworthy, on the exact axis this category is decided. If that is interesting, I would love to talk.

— Vishal · IDYeah Studio

What this is grounded in

This note is grounded in the material you shared and what I could reach myself: the MOSAIX US deck, and 24 screenshots of the live demo instance (an Org Admin session, no live customer data, trial expiring 18 July 2026), plus some 2026 category and buyer research to check my read against the market. Screens covered include onboarding (workspace setup, trial start, the sample-org welcome), Home, My Feed, Skill Assessments, My Career Arc, Career Recommendations, Project Marketplace, Mentorverse, My Community, My Team, Organization Design Engine, Talent Mapper, Succession Planning, Succession Analytics, Talent Metrics, Skills Framework Builder, HR Operations, and the four “Coming Soon” modules.

Market and buyer context from 2026 talent-intelligence analyses (Knowlee, Engagedly), enterprise HR buying guides (UKG, SelectSoftwareReviews) and independent HR-AI analysis (Viewpoint Analysis).

Prepared by IDYeah Studio · Product Design & Brand Storytelling · July 2026. An independent design read, made to inform a possible redesign. Confidential.

Researched and pressure-tested with **Crit**, IDYeah’s in-house audit method. The read, and the calls, are mine.